

# Machine learning memo

Daniel Perez

## 1 Model representation

### 1.1 Symbols

- $m$  = number of training example
- $x$ 's = "input" variable / features
- $y$ 's = "output" variable / "target" variable
- $(x, y)$  - one training example
- $(x^{(i)}, y^{(i)})$  -  $i^{\text{th}}$  training example
- $h : x \rightarrow y$  - hypothesis function (takes input and estimates output)

#### 1.1.1 Linear regression

Also called univariate linear regression.

$h$  is a linear function.

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

### 1.2 Cost function

$\theta_i$  are the parameters of the equation. For linear regression:  $h_{\theta}(x) = \theta_0 + \theta_1 x$

Cost function (square error cost function)

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

We want to minimize over the parameters  $\theta_0$  and  $\theta_1$ :

$$\min_{\theta_0, \theta_1} J(\theta_0, \theta_1)$$

The square error cost function is one of the most used for regression.

## 2 Multiple features

### 2.1 Notations

- $n$  is the number of features
- $x^{(i)}$ : input features of  $i^{\text{th}}$  training example
- $x_j^{(i)}$ : input feature  $j$  of  $i^{\text{th}}$  training example

## 2.2 Hypothesis

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

For convenience,  $x_0 = 1$ , so that

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \in \mathbb{R}^{n+1} \quad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

we can therefore write

$$h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n \quad (1)$$

$$= \theta^T x \quad (2)$$

which means

$$h_{\theta}(x) = \underbrace{[\theta_0 \quad \theta_1 \quad \cdots \quad \theta_n]}_{\theta^T} \cdot \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix}$$

This is also called Multivariate linear regression.

## 2.3 Gradient descent

- Parameters:  $\theta_0, \theta_1, \dots, \theta_n$  which we think of as  $\theta \in \mathbb{R}^{n+1}$
- Cost function:  $J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$  where  $\theta = [\theta_0 \quad \theta_1 \quad \cdots \quad \theta_n]$

Gradient descent: repeat

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

until we converge.

with simultaneous update. By computing the partial derivative, this gives us

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

## 2.4 Feature scaling

Use features on the same scale (with same range of values).

e.g.

- Size of the bedroom (100 - 2000feet<sup>2</sup>)
- Number of beds (1 - 5)

We will want to have values such as  $-1 \leq x_i \leq 1$  Otherwise, gradient descent will converge slowly.

## 2.5 Mean normalization

Replace  $x_i$  by  $x_i - \mu_i$ , so that the mean of each feature is around 0

Combining both:

$$x_i \leftarrow \frac{x_i - \mu_i}{s_i}$$

where  $s_i$  can be either the the range (max - min) or the standard derivation of the feature.

## 2.6 Learning rate

Gradient descent should decrease after each iteration.

Plotting cost function ( $y$  axis) and number of iteration ( $x$  axis) helps to make sure gradient descent is working.

If cost function increases with gradient descent increases, use smaller learning rate  $\alpha$ . Same if the cost goes up and down.

- For sufficiently small  $\alpha$ ,  $J(\theta)$  should decrease on every iteration
- But if  $\alpha$  is too small, gradient descent can be too slow

When choosing  $\alpha$ , try a range of values:  $\dots, 0.001, 0.003, 0.01, 0.03, 0.1, 0.3, 1, \dots$

## 2.7 Normal equation

Given  $X$  is the matrix containing features, with  $x_0 = 1$  and  $y$  is a vector containing the results,

$$\theta = (X^T X)^{-1} X^T y$$

In details:

$$x^{(i)} = \begin{bmatrix} x_0^{(i)} \\ x_1^{(i)} \\ x_2^{(i)} \\ \vdots \\ x_n^{(i)} \end{bmatrix} \quad X = \begin{bmatrix} \dots (x^{(1)})^T \dots \\ \dots (x^{(2)})^T \dots \\ \vdots \\ \dots (x^{(m)})^T \dots \end{bmatrix}$$

$X$  will have a dimension of  $m \times (n + 1)$

## 3 Logistic regression

Algorithm to classify a data set in different categories.

### 3.1 Hypothesis

$$h_\theta(x) = g(\theta^T x) \tag{3}$$

$$g(z) = \frac{1}{1 + e^{-z}} \tag{4}$$

$g(z)$  has the following property:

$$g(z) \geq 0.5 \leftrightarrow z \geq 0$$

### 3.2 Cost function

We want the cost function to be convex, so we cannot use the same function as for linear regression.

We define the cost function as follow:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x), y)$$

For linear regression, the  $\text{Cost}(h_{\theta}(x), y)$  was the squared error.

For logistic regression, we use

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

More intuitively, if we predicted with certainty  $y = 0$  but it turned out that  $y = 1$ , we pay a very large cost.

On the opposite, if we predicted with certainty that  $y = 0$  and it was correct, the cost is very close to 0.

The above function can be rewritten as follow:

$$\text{Cost}(h_{\theta}(x), y) = -y \log(h_{\theta}(x)) - (1 - y) \log(1 - h_{\theta}(x))$$

we therefore get

$$J(\theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right]$$

we then need to compute  $\min_{\theta} J(\theta)$ , which can be achieved by using the gradient descent.

We compute the partial derivative as follow

$$\frac{\partial}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

### 3.3 Optimizations

There are more optimized algorithms than gradient descent that can be used to compute the parameters. For example:

- Conjugate gradient
- BFGS
- L-BFGS

Advantages:

- No need to manually pick  $\alpha$
- Often faster than gradient descent

Disadvantages

- More complex

These can be used easily in octave as follow. Given the following example,

$$\theta = \begin{bmatrix} \theta_1 \\ \theta_1 \end{bmatrix} \quad (5)$$

$$J(\theta) = (\theta_1 - 5)^2 + (\theta_2 - 5)^2 \quad (6)$$

$$\frac{\partial}{\partial \theta_1} J(\theta) = 2(\theta_1 - 5) \quad (7)$$

$$\frac{\partial}{\partial \theta_2} J(\theta) = 2(\theta_2 - 5) \quad (8)$$

```
function [jVal, gradient] = costFunction(theta)
    jVal = (theta(1) - 5)^2 + (theta(2) - 5)^2;
    gradient = zeros(2, 1);
    gradient(1) = 2 * (theta(1) - 5);
    gradient(2) = 2 * (theta(2) - 5);
end

options = optimset('GradObj', 'on', 'MaxIter', 100);
initialTheta = zeros(2, 1);
[optTheta, functionVal, exitFlag] = fminunc(@costFunction, initialTheta, options);
```

### 3.4 Multiclass classification

We can use the one-vs-all algorithm.

Train a logistic regression classifier  $h_{\theta}^{(i)}(x)$  for each class  $i$  to predict the probability that  $y = i$   
 On a new input  $x$ , to make a prediction, pick the class  $i$  that maximizes

$$\max_i h_{\theta}^{(i)}(x)$$

## 4 Regularization

### 4.1 Overfitting

If we have too many features, the learned hypothesis may fit the training set very well ( $J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 \approx 0$ ), but fail to generalize to new examples.

#### 4.1.1 Addressing overfitting

Options:

1. Reduce number of features
  - Manually select which features to keep
  - Model selection algorithm
2. Regularization
  - Keep all the features, but reduce the magnitude/values of parameters  $\theta_j$
  - Works well when we have a lot of features, each of which contributes a bit to predicting  $y$

## 4.2 Cost function

Small values for parameters  $\theta_0, \theta_1, \dots, \theta_n$

- Simpler hypothesis
- Less prone to overfitting

To implement it, we can change the cost function as follow

$$J(\theta) = \frac{1}{2m} \left[ \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2 + \lambda \sum_{j=1}^m \theta_j^2 \right]$$

$\lambda$  is called the regularization parameter.

## 4.3 Regularized linear regression

With

$$J(\theta) = \frac{1}{2m} \left[ \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2 + \lambda \sum_{j=1}^m \theta_j^2 \right]$$

we want to compute  $\min_{\theta} J(\theta)$ .

We have

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right) x_j^{(i)} + \frac{\lambda}{m} \theta_j$$

therefore we only need to separate the computation of  $\theta_0$  and  $\theta_j$  ( $j \geq 1$ ) so we do not regularize  $\theta_0$ .

### 4.3.1 Gradient descent

The gradient descent update for  $\theta_j$  looks like

$$\theta_j = \theta_j - \alpha \left[ \frac{1}{m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right) x_j^{(i)} + \frac{\lambda}{m} \theta_j \right] \quad (9)$$

$$= \theta_j \left( 1 - \alpha \frac{\lambda}{m} \right) - \alpha \frac{1}{m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right) x_j^{(i)} \quad (10)$$

We usually want  $(1 - \alpha \frac{\lambda}{m}) < 1$

### 4.3.2 Normal equation

Given

$$X = \begin{bmatrix} (x^{(1)})^T \\ \vdots \\ (x^{(m)})^T \end{bmatrix} \quad y = \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(m)} \end{bmatrix}$$

we are looking for  $\min_{\theta} J(\theta)$

The normal equation in this case is given by

$$\theta = \left( X^T X + \lambda \begin{bmatrix} 0 & & & \\ & 1 & & \\ & & \ddots & \\ & & & 1 \end{bmatrix} \right)^{-1} X^T y$$

#### 4.4 Regularized logistic regression

We update our cost function to look as follow:

$$J(\theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

To implement the gradient descent, we need to use the same steps as for linear regression.

### 5 Neural networks

Neural network is made of

1. An input layer
2. 0 or more hidden layers
3. An output layer

#### 5.1 Example and computations

Given a neural network with 3 units in the input layer, and a single hidden layer, the computations are defined as follow:

$$a_1^{(2)} = g \left( \Theta_{10}^{(1)} x_0 + \Theta_{11}^{(1)} x_1 + \Theta_{12}^{(1)} x_2 + \Theta_{13}^{(1)} x_3 \right) \quad (11)$$

$$a_2^{(2)} = g \left( \Theta_{20}^{(1)} x_0 + \Theta_{21}^{(1)} x_1 + \Theta_{22}^{(1)} x_2 + \Theta_{23}^{(1)} x_3 \right) \quad (12)$$

$$a_3^{(2)} = g \left( \Theta_{30}^{(1)} x_0 + \Theta_{31}^{(1)} x_1 + \Theta_{32}^{(1)} x_2 + \Theta_{33}^{(1)} x_3 \right) \quad (13)$$

$$h_{\Theta}(x) = g \left( \Theta_{10}^{(2)} a_0^{(2)} + \Theta_{11}^{(2)} a_1^{(2)} + \Theta_{12}^{(2)} a_2^{(2)} + \Theta_{13}^{(2)} a_3^{(2)} \right) \quad (14)$$

We define

$$\Theta_{10}^{(1)} x_0 + \Theta_{11}^{(1)} x_1 + \Theta_{12}^{(1)} x_2 + \Theta_{13}^{(1)} x_3 = z_1^{(2)}$$

so that

$$a_1^{(2)} = g(z_1^{(2)})$$

Given the above example, we have

$$x = a^{(1)} = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{bmatrix} \quad z^{(2)} = \begin{bmatrix} z_1^{(2)} \\ z_2^{(2)} \\ z_3^{(2)} \end{bmatrix}$$

and we can use the following computations.

$$z^{(2)} = \Theta^{(1)} a^{(1)} \quad (15)$$

$$a^{(2)} = g\left(z^{(2)}\right) \quad (16)$$

To compute the next layer (output layer for this example), add  $a_0^{(2)} = 1$ , and repeat:

$$z^{(3)} = \Theta^{(2)} a^{(2)} \quad (17)$$

$$h_{\Theta}(x) = a^{(3)} = g\left(z^{(3)}\right) \quad (18)$$

## 5.2 Cost function

- $\left(x^{(1)}, y^{(1)}\right), \left(x^{(2)}, y^{(2)}\right), \dots, \left(x^{(m)}, y^{(m)}\right)$ : training set
- $L$ : total number of layers in network
- $s_l$ : number of units (not counting bias unit) in layer  $l$

$$J(\Theta) = -\frac{1}{m} \left[ \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log \left( h_{\Theta}(x^{(i)}) \right)_k + (1 - y_k^{(i)}) \log \left( 1 - (h_{\Theta}(x^{(i)}))_k \right) \right] \quad (19)$$

$$+ \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} \left( \Theta_{ji}^{(l)} \right) \quad (20)$$

## 5.3 Back propagation algorithm

$\delta_j^{(l)}$  is the “error” of node  $j$  in layer  $l$ .

For each output unit (here layer  $L = 4$ )

$$\delta_j^{(4)} = a_j^{(4)} - y_j \quad (21)$$

$$\delta_j^{(3)} = \left( \Theta^{(3)} \right)^T \delta^{(4)} \star g' \left( z^{(3)} \right) \quad (22)$$

$$\delta_j^{(2)} = \left( \Theta^{(2)} \right)^T \delta^{(3)} \star g' \left( z^{(2)} \right) \quad (23)$$

where

$$g' \left( z^{(3)} \right) = a^{(3)} \star \left( 1 - a^{(3)} \right)$$

and

$$\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) = a_j^{(l)} \delta_i^{(l+1)} \quad \text{if } \lambda = 0$$

### 5.3.1 Algorithm

- Set  $\Delta_{ij}^{(l)} = 0$  (for all  $l, i, j$ )
- For  $i = 1$  to  $m$ 
  - Set  $a^{(1)} = x^{(i)}$



- Perform forward propagation to compute  $a^{(l)}$  for  $l = 2, 3, \dots, L$
- Using  $y^{(i)}$ , compute  $\delta^{(L)} = a^{(L)} - y^{(i)}$
- Compute  $\delta^{(L-1)}, \delta^{(L-2)}, \dots, \delta^{(2)}$ 

$$\Delta_{ij}^{(l)} = \Delta_{ij}^{(l)} + a_j^{(l)} \delta_i^{(l+1)}$$
- $D_{ij}^{(i)} = \frac{1}{m} \Delta_{ij}^{(l)} + \lambda \Theta_{ij}^{(l)}$  if  $j \neq 0$
- $D_{ij}^{(i)} = \frac{1}{m} \Delta_{ij}^{(l)}$  if  $j = 0$

## 6 Evaluating algorithms

Precision:

$$\frac{\text{\#true positives}}{\text{\#true positives} + \text{\#false positives}}$$

Recall:

$$\frac{\text{\#true positives}}{\text{\#true positives} + \text{\#false negatives}}$$

$F_1$  score:

$$2 \frac{PR}{P + R}$$